

LECTURER: Anna Androvitsanea

DATA ENGINEERING

Database Engineering — Unit 1

Dr.-Ing Anna Androvitsanea

IU Internationale Hochschule GmbH

January 29, 2024

Foundation of Data Systems	1
Data Processing at Scale	2
Microservices	3
Governance & Security	4
Common Cloud Platforms & Services	5
DataOps	6

Introductory Round

- ▶ Who are you?
- ▶ Anna
- ▶ iu / Flightright GmbH
- ▶ Data Scientist and ML Engineer
- ▶ Have worked as a Civil Engineer and Archaeologist
- ▶ Data Scientist enthusiast

Role of Data Engineers and Data Scientists

- ▶ Exploring the distinct roles and responsibilities of data engineers and data scientists.
- ▶ Examining the interplay and collaboration between these two roles in a real-world project.
- ▶ Case Study: A successful collaboration between data engineering and data science teams in a retail analytics project.

Case Study Overview: Retail Analytics Project

- ▶ Introduction to the retail analytics project.
- ▶ Objectives and challenges of the project.
- ▶ Overview of the teams involved: Data Engineering and Data Science.

Role of the Data Engineering Team

- ▶ Responsibilities and tasks of the data engineering team in the project.
- ▶ Techniques and technologies used for data collection, storage, and preprocessing.
- ▶ How the team ensured data quality and integrity.

Role of the Data Science Team

- ▶ The data science team's approach to data analysis and modeling.
- ▶ Specific data science methods and tools used.
- ▶ Insights and predictions derived from the data.

Collaboration Between Teams

- ▶ How the teams worked together: communication and workflow strategies.
- ▶ Integration of data engineering outputs with data science models.
- ▶ Addressing challenges and adapting strategies in the collaborative process.

Outcomes and Achievements

- ▶ Key results and successes of the project.
- ▶ Impact of the collaboration on project outcomes.
- ▶ Lessons learned and best practices identified for future projects.

Unit 1

FOUNDATIONS OF DATA SYSTEMS

Study Goals

- ▶ Explain properties of well-designed data-intensive systems.
- ▶ Learn how to quantify the reliability of a system.
- ▶ Understand how to compare different approaches to system scalability.
- ▶ Discuss aspects of maintainability in the development of data-intensive systems.

EXPLAIN SIMPLY

1. List **three aspects** that enable modern data-intensive data systems to process large amounts of data while running stably.
2. Explain ways to **adapt** these systems to higher or lower **performance requirements**.
3. Discuss how these systems can be designed for straightforward operation in terms of **maintenance and evolution**.

RECAP OF THE UNIT'S KEY TOPICS

Reliability

- Reliability Metrics
- Hardware Fault
- Software Error
- Human Error

Scalability

- Load and Throughput
- Service Level Agreements (SLAs)
- Scale Up – Scale Out

Maintainability

- Operability
- Simplicity
- Evolvability

Source of the text: Hartz, Walker & Mahar, 1997.

Properties of Well-Designed Data-Intensive Systems: Reliability

- ▶ Reliability: Ability to function correctly even in the face of adversity.
- ▶ Includes handling user errors, system faults, and environmental challenges.

Scalability

- ▶ Scalability: Capacity to handle growth and increased load effectively.
- ▶ Involves scaling up/out and managing load distribution efficiently.

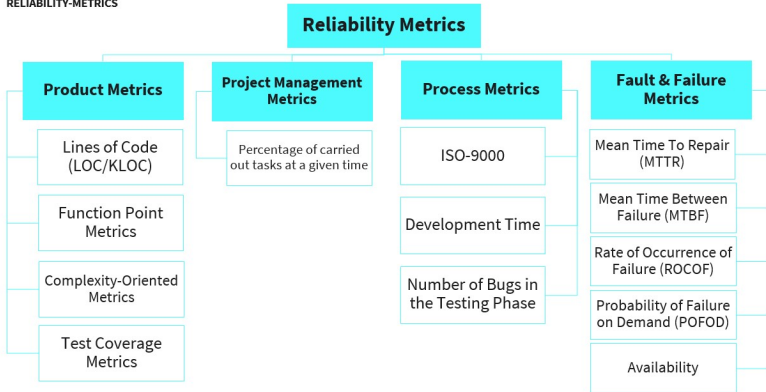
Maintainability

- ▶ Maintainability: Ease of maintaining system operations and updates.
- ▶ Focuses on simplicity, good documentation, and automation of operational tasks.

Properties of Well-Designed Data-Intensive Systems

Reliability

RELIABILITY-METRICS



Source of the text: Wilkins, 2002.

Source of the image: Müller-Kett (2022) based on Hartz et al., 1997.

Reliability Metric: Lines of Code (LOC/KLOC)

- ▶ **Explanation:** LOC (Lines of Code) and KLOC (Thousand Lines of Code) measure software size or complexity.
- ▶ **Example:** Using LOC to estimate effort or complexity of software projects. Larger projects typically have more LOC.

Reliability Metric: Function Point Metrics

- ▶ **Explanation:** Measures software functionality from the user's perspective.
- ▶ **Example:** Used in budgeting and planning, assessing the functionality delivered to users.

Project Management Metric: Development Time

- ▶ **Explanation:** Tracks the total time taken from project initiation to completion.
- ▶ **Example:** Used to assess project efficiency, schedule adherence, and predict future project timelines.

Process Metric: Number of Bugs in Testing Phase

- ▶ **Explanation:** Counts the number of defects found during the testing phase of software development.
- ▶ **Example:** Indicator of software quality before release; higher numbers may indicate issues in earlier development stages.

Fault Failure Metric: Mean Time To Repair (MTTR)

- ▶ **Explanation:** Measures the average time taken to repair a fault or failure in the system.
- ▶ **Example:** Used in maintenance and reliability analysis; shorter MTTR implies better system maintainability.

Fault Failure Metric: Mean Time Between Failures (MTBF)

- ▶ **Explanation:** Indicates the average time between failures in a system.
- ▶ **Example:** Used to gauge the reliability of a system; longer MTBF is desirable.

Fault Failure Metric: Rate of Occurrence of Failure (ROCOF)

- ▶ **Explanation:** Measures the frequency of failures in a system.
- ▶ **Example:** Useful for identifying trends in system reliability over time.

Fault Failure Metric: Probability of Failure on Demand (POFOD)

- ▶ **Explanation:** Assesses the likelihood of a system failing when required.
- ▶ **Example:** Critical for systems where reliability on demand is essential, like safety-critical systems.

Fault Failure Metric: Availability

- ▶ **Explanation:** Measures the proportion of time a system is operational and available for use.
- ▶ **Example:** High availability is crucial for systems requiring constant uptime, like web servers.

HARDWARE FAULT


Physical damage in one machine's components

- Redundancies (i.e., RAIDs)
- Redundancy across multiple machines

Power supply blackouts /

Network disruptions

- Redundancies across multiple data centers/availability zones



Source of the text: Deltienne, 2002.
Source of the image : Microsoft Archive.

Hardware Faults

- ▶ Physical damage in one machine's components.
- ▶ Strategies to mitigate:
 - ▶ Implementing redundancies, like RAID systems.
 - ▶ Redundancy across multiple machines.

Understanding RAID Systems

- ▶ RAID (Redundant Array of Independent Disks) is a data storage technology.
- ▶ Combines multiple physical disk drive components into one or more logical units.
- ▶ Key Purposes:
 - ▶ Data redundancy: Protects data by duplicating it across multiple disks.
 - ▶ Improved performance: Can increase read and write speeds by distributing operations across several disks.
- ▶ Common RAID Levels:
 - ▶ RAID 0: Striped set without parity or mirroring.
 - ▶ RAID 1: Mirroring without parity or striping.
 - ▶ RAID 5: Block-level striping with distributed parity.

Redundancies


- ▶ Essential for mitigating hardware and network failures.
- ▶ Types of Redundancies:
 - ▶ In-device redundancies (like RAIDs).
 - ▶ Across multiple devices or systems.
 - ▶ Redundancies across data centers or availability zones.

Network Disruptions

- ▶ Causes and consequences of network disruptions.
- ▶ Mitigation strategies:
 - ▶ Redundancies across multiple data centers.
 - ▶ Diverse network routes to prevent single points of failure.

SOFTWARE ERROR

- Memory leakage
- Service stops working
- Cascading failure
- Software is a mental product and humans will introduce error to this



Source of the image: Naylor & Joyner, 2014.
Source of the image: Microsoft Archive.

MENSCHLICHER ERROR

- Humans might introduce considerable potential for error into systems
- 32% of security incidents are caused by employees
- Provide efficient recovery tools
- Provide minimum opportunity for human error
- Intensive testing



Source of the text: BakerHostetler, 2017.
Source of the image: Microsoft Archive.

LOAD & THROUGHPUT

- **System Load**
 - Read-to-Write ratio
 - Number of online users
 - Requests per time
- **System Performance Load**
 - Running times for read/write operations
 - Throughput
 - I/O-Rate: Input-Output Operations per Second (IOPS)
 - Response Time

Properties of Well-Designed Data-Intensive Systems

Scalability

System Load

- ▶ System Load refers to the volume of work handled by a system.
- ▶ Key Metrics:
 - ▶ Read-to-Write ratio: Balance of reading vs. writing operations.
 - ▶ Number of online users: Concurrent users interacting with the system.
 - ▶ Requests per time: Frequency of operations requested from the system.

System Performance Load


- ▶ Focuses on the efficiency of processing operations.
- ▶ Measures:
 - ▶ Running times for read/write operations.
 - ▶ How quickly the system can process tasks.

Throughput

- ▶ Throughput: Quantity of work a system can handle in a given time.
- ▶ Key Indicators:
 - ▶ I/O-Rate (Input-Output Operations per Second, IOPS).
 - ▶ Response Time: The time it takes for the system to respond to a request.

SCALE UP - SCALE OUT

- Optimize the **software-side**
- Scale the **physical ressources**



Source of the image: Alvarid (2022).

Scale Up: Optimizing Software and Resources

- ▶ **Definition:** Scale Up (Vertical Scaling) involves increasing the power of existing hardware or software.
- ▶ **Optimize the Software-Side:**
 - ▶ Enhancing the performance of the application to better utilize existing resources.
 - ▶ Examples: optimizing algorithms, refining code for efficiency.
- ▶ **Scale the Physical Resources:**
 - ▶ Upgrading existing hardware capabilities: more RAM, faster CPU, larger storage.
 - ▶ Impact: Improved performance without changing the system architecture.

Scale Out: Expanding Capacity Horizontally

- ▶ **Definition:** Scale Out (Horizontal Scaling) involves adding more nodes to a system to distribute the load.
- ▶ **Optimize the Software-Side:**
 - ▶ Designing software to efficiently run on multiple servers or nodes.
 - ▶ Examples: Implementing load balancing, ensuring data consistency across nodes.
- ▶ **Scale the Physical Resources:**
 - ▶ Adding more machines or instances to the existing pool.
 - ▶ Impact: Enhances capacity and reliability, allows for handling increased load without a single point of failure.

Properties of Well-Designed Data-Intensive Systems


Maintainability

MAINTAINABILITY

Making it easy to install, use and modify the system

Minimizing the discomfort of legacy maintenance

- Corrective maintenance
- Adaptive maintenance
- Perfective maintenance
- Preventive maintenance




Source of the image: Müller-Kett (2022).

OPERABILITY

Operators are responsible for

- Monitoring the system's health
- Installing security patches
- Predicting and solving potential problems
- Ensuring a stable production environment
- Preserving the knowledge about the system




Source of the text: Hamilton, 2007, p. 231-242.
Source of the image: Müller-Kett (2022).

OPERABILITY

Making it easy for operators to use the system

- Providing transparent monitoring possibilities
- Supporting automation and integration with standard tools
- Providing concise documentation
- Ensuring service/machine independence
- Providing default setups and self-healing features




Source of the text: Hamilton, 2007, p. 231-242.
Source of the image: Müller-Kett (2022).

SIMPLICITY

Making the code as easy as possible
and only as complex as it is necessary

Reducing...

- accidental complexity
 - long code
 - module interactions & dependencies
 - inconsistent naming
-
- increasing the level of abstraction



Source of the image: Müller-Kett (2022).

Simplifying Code

- ▶ Aim: Make code as easy as possible, only complex as necessary.
- ▶ Benefits: Easier to understand, maintain, and less prone to errors.

Reducing Accidental Complexity

- ▶ Accidental Complexity: Unnecessary complexity not inherent to the problem.
- ▶ Strategies: Refactoring code, using design patterns, simplifying logic.

Streamlining Code Length and Modules

- ▶ Reducing Long Code: Break down into smaller, manageable functions.
- ▶ Managing Module Interactions: Minimize dependencies, use clear interfaces.

Consistent Naming and Abstraction

- ▶ Consistent Naming: Improves readability and maintainability.
- ▶ Increasing Level of Abstraction: Focus on the 'what' over 'how'.

EVOLVABILITY

Making changes to the system as easy as possible

- Similar to the term “agile”
- DataOps



Source of the image: Müller-Kett (2022).

Evolvability in Data Systems

- ▶ Objective: Making changes to the system as easy as possible.
- ▶ Similar to the agile methodology: Adaptability and responsiveness to change.
- ▶ Example: Iterative development, where systems are continuously improved.

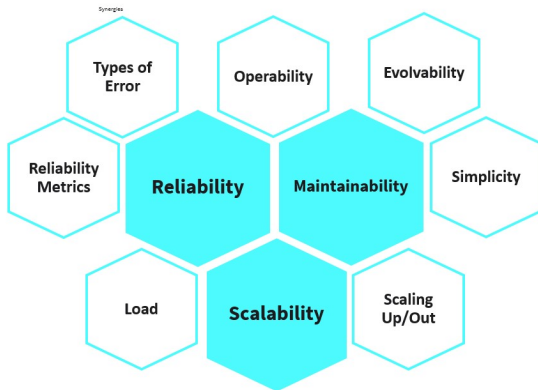
DataOps: Facilitating Evolvability

- ▶ DataOps: Agile and lean practices in data management.
- ▶ Example: Automating data pipelines to quickly adapt to new data sources.
- ▶ Benefits: Faster deployment, improved quality, and more efficient data use.

Example of a Data Pipeline

- ▶ **Objective:** Move and process data from source to destination.
- ▶ **Stages:**
 1. *Extraction:* Collect data from various sources (e.g., databases, online services).
 2. *Transformation:* Process data (e.g., cleaning, aggregating, formatting).
 3. *Loading:* Load transformed data into a destination (e.g., data warehouse).
- ▶ **Use Case:**
 - ▶ Extract sales data from online transactions.
 - ▶ Aggregate sales by region and calculate total revenue.
 - ▶ Load processed data into a data warehouse for analysis.

SYNERGIES



Source of the image: Müller-Kett (2022).

Study Goals

- ▶ Explain properties of well-designed data-intensive systems.
- ▶ Learn how to quantify the reliability of a system.
- ▶ Understand how to compare different approaches to system scalability.
- ▶ Discuss aspects of maintainability in the development of data-intensive systems.

Study Goals: Key Aspects of Modern Data-Intensive Systems

- ▶ **Distributed Computing**
 - ▶ Utilizes multiple machines for efficient data handling.
 - ▶ Improves resilience and scalability.
- ▶ **Advanced Data Storage Solutions**
 - ▶ Incorporates technologies like NoSQL databases.
 - ▶ Facilitates efficient data management and retrieval.
- ▶ **High-Performance Processing Frameworks**
 - ▶ Uses frameworks like Hadoop and Spark.
 - ▶ Optimized for rapid analysis and processing of big data.

Study Goals: Quantifying System Reliability

- ▶ Understanding Key Reliability Metrics:
 - ▶ Mean Time Between Failures (MTBF).
 - ▶ Mean Time To Repair (MTTR).
 - ▶ System Availability and Uptime.
- ▶ Applying Reliability Metrics:
 - ▶ Calculating reliability scores.
 - ▶ Analyzing failure rates and recovery processes.
- ▶ Case Study: Examining real-world systems' reliability reports.

Study Goals: Comparing Approaches to System Scalability

- ▶ Understanding Scalability:
 - ▶ Definition and importance in data-intensive systems.
- ▶ Scalability Approaches:
 - ▶ Vertical Scaling (Scaling Up): Increasing the power of existing hardware.
 - ▶ Horizontal Scaling (Scaling Out): Adding more nodes to a system.
- ▶ Evaluating Scalability Strategies:
 - ▶ Cost-effectiveness, performance implications, and maintenance considerations.
 - ▶ Real-world examples and case studies for context and understanding.

Study Goals: Maintainability in Data-Intensive Systems

- ▶ Importance of Maintainability:
 - ▶ Key to long-term success and adaptability of systems.
- ▶ Aspects of Maintainability:
 - ▶ Code Simplicity and Clarity: Easy to understand and modify.
 - ▶ Documentation Quality: Comprehensive and up-to-date.
 - ▶ Automated Testing: Ensures reliability and ease of changes.
 - ▶ Modular Design: Allows independent development and updates.
- ▶ Strategies for Improvement:
 - ▶ Regular code reviews, refactoring, and keeping up with technological advancements.

EXPLAIN SIMPLY

1. List **three aspects** that enable modern data-intensive data systems to process large amounts of data while running stably.
2. Explain ways to **adapt** these systems to higher or lower **performance requirements**.
3. Discuss how these systems can be designed for straightforward operation in terms of **maintenance and evolution**.

EXPLAIN SIMPLY: Key Aspects for Processing Large Data Volumes

- ▶ **Distributed Computing:** Utilizing multiple machines for data processing.
- ▶ **Advanced Storage Systems:** Implementing scalable and efficient storage solutions like NoSQL.
- ▶ **High-Performance Processing:** Leveraging frameworks like Hadoop or Spark for fast data processing.

EXPLAIN SIMPLY: Adapting Systems to Performance Requirements

- ▶ **Scalability:** Vertical scaling (enhancing existing systems) and horizontal scaling (adding more systems).
- ▶ **Resource Optimization:** Tailoring resource allocation based on current load and performance targets.
- ▶ **Performance Tuning:** Adjusting configurations for optimal efficiency under varying workloads.

EXPLAIN SIMPLY: Designing for Maintenance and Evolution

- ▶ **Modular Architecture:** Facilitates updates and integration of new technologies.
- ▶ **Automation and Testing:** Ensures reliability and eases updates.
- ▶ **Documentation and Standardization:** Provides clarity and consistency for future development.

SESSION 1

TRANSFER TASK

TRANSFER TASK

A start-up that sells **sustainable products in smaller stores** has been very successful in recent years. As a result, more stores are to be opened worldwide. As a Data Engineer, you will be tasked with **designing the data system that stores and processes data** about the products offered and their suppliers.

As a team, develop **key points** to ensure that this system will run appropriately effectively and perform well. For each of these points, think about **specific measures to be implemented** in the system.

TRANSFER TASK
PRESENTATION OF THE RESULTS

Please present your
results.

The results will be
discussed in plenary.



- **Reliability metrics** are implemented in the system
 - To ensure a fast implementation, the following metrics are implemented first (based on log file monitoring) and monitored in a dashboard, for example
 - Product metrics (lines of code, function point metrics, test coverage metrics)
 - Process metrics (number of bugs in the test phase)
 - Fault & Failure metrics (Availability, Rate of Occurrence of Failure (ROCOF), Mean Time To Repair (MTTR))
- To keep the system functional in case of **hardware failures**, it is implemented on multiple redundant computing units
 - Virtualization in containers
 - Distribution of the system over several data centers
 - Distribution of the system over several geographically separated "Availability Zones"

- The system is implemented in **collaboration** with an established **cloud provider**
 - This way, we ensure that the system is dynamically scalable at all times
- We use both **vertical and horizontal scaling** to adapt the system to the system load
- To maintain a defined **service level**, the system is constantly monitored for load and throughput based on log files
- Scale Up - Scale Out

We ensure the **operability** of the system by....

- making the system **easy to monitor** in the form of a dashboard and automated alert mechanisms
- **simplifying deployment** through a predefined CI/CD pipeline
- providing **clear documentation**
- **isolate** individual **services** from each other
- implementing predefined **default settings** that can be overridden if necessary

We keep the system **simple** by...

- regularly checking if the code can be **simplified**
- keeping the code as **short** as possible
- **avoiding interactions** between modules
- using **abstraction** (e.g. by self-contained functions and modules)
- set up rules for **consistent naming**

In addition, we design the system in such a way that it can be **easily modified or extended**

- for example, we implement all self-contained functionalities as **independent functions or modules**



1. Which task does not fall under the responsibility of data engineers?

- a) Data Analysis
- b) Removing corrupted data
- c) Acquiring data from various sources
- d) Optimizing databases for analysis



2. Which of the following choices is the term used to describe the probability of error-free software operation over a given period of time in a given environment?
- a) Reliability
 - b) Maintainability
 - c) Scalability
 - d) Security



3. If some errors are found in the application and they need to be fixed, what kind of maintenance is used?
- a) adaptive maintenance
 - b) perfecting maintenance
 - c) preventive maintenance
 - d) corrective maintenance

LIST OF SOURCES

Text:

BakerHostetler. (2017, May 9). *Webinar: Be compromise ready: Go back to the basics*. <https://www.bakerlaw.com/events/webinar-be-compromise-ready-go-back-to-thebasics>

Détienne, F. (2002). *Software design: Cognitive aspects*. Springer.

Hamilton, J. (2007). On designing and deploying internet-scale services. *Proceedings of the 21st Large Installation System Administration Conference (LISA '07)*, 231–242.

Hartz, M. A., Walker, E. L. & Mahar, D. (1997). *Introduction to software reliability: A state of the art review*. The Center.

Naylor, W., & Joyner, B. (2014). A discourse on software safety and software reliability. *2014 Reliability and Maintainability Symposium*. <https://ieeexplore.ieee.org/document/6798493>

Wilkins, D. J. (2002). The Bathtub Curve and Product Failure Behavior (Part 1 of 2). *Reliability HotWire*, 21. Abgerufen am 15. März 2022, von <https://www.weibull.com/hotwire/issue21/hottopics21.htm>

Image:

Alavirad, 2020.

Müller-Kett (2022).

Müller-Kett (2022) based on Hartz et al., 1997.

Microsoft Archive.